

Machine Learning - Explain models with SHAP

Classifier

<https://www.kaggle.com/parulpandey/palmer-archipelago-antarctica-penguin-data>

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import plot_tree
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix, plot_confusion_matrix
import shap
```

```
In [ ]: penguins = pd.read_csv('penguins_size.csv')
penguins.head()
```

```
Out[ ]:   species  island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex
0  Adelie  Torgersen         39.1           18.7           181.0         3750.0  MALE
1  Adelie  Torgersen         39.5           17.4           186.0         3800.0  FEMALE
2  Adelie  Torgersen         40.3           18.0           195.0         3250.0  FEMALE
3  Adelie  Torgersen         NaN           NaN           NaN           NaN     NaN
4  Adelie  Torgersen         36.7           19.3           193.0         3450.0  FEMALE
```

```
In [ ]: penguins.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column              Non-Null Count  Dtype
---  -
0   species             344 non-null   object
1   island              344 non-null   object
2   culmen_length_mm    342 non-null   float64
3   culmen_depth_mm     342 non-null   float64
4   flipper_length_mm   342 non-null   float64
5   body_mass_g         342 non-null   float64
6   sex                 334 non-null   object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

```
In [ ]: penguins.describe(include='all')
```

```
Out[ ]:   species  island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex
count      344      344           342.000000           342.000000           342.000000           342.000000  334
unique         3         3              NaN              NaN              NaN              NaN         3
top    Adelie  Biscoe              NaN              NaN              NaN              NaN  MALE
freq     152     168              NaN              NaN              NaN              NaN  168
mean     NaN     NaN           43.921930           17.151170           200.915205          4201.754386  NaN
std     NaN     NaN           5.459584           1.974793           14.061714           801.954536  NaN
min     NaN     NaN           32.100000           13.100000           172.000000          2700.000000  NaN
25%     NaN     NaN           39.225000           15.600000           190.000000          3550.000000  NaN
50%     NaN     NaN           44.450000           17.300000           197.000000          4050.000000  NaN
75%     NaN     NaN           48.500000           18.700000           213.000000          4750.000000  NaN
max     NaN     NaN           59.600000           21.500000           231.000000          6300.000000  NaN
```

```
In [ ]: penguins.isna().sum()
```

```
Out[ ]: species      0
island      0
culmen_length_mm    2
culmen_depth_mm     2
```

```
flipper_length_mm    2
body_mass_g          2
sex                  10
dtype: int64
```

```
In [ ]: penguins = penguins.dropna()
penguins = penguins[penguins['sex'] != '.']
penguins = penguins.reset_index(drop=True)
penguins.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 333 entries, 0 to 332
Data columns (total 7 columns):
#   Column              Non-Null Count  Dtype
---  -
0   species              333 non-null   object
1   island               333 non-null   object
2   culmen_length_mm     333 non-null   float64
3   culmen_depth_mm     333 non-null   float64
4   flipper_length_mm   333 non-null   float64
5   body_mass_g         333 non-null   float64
6   sex                  333 non-null   object
dtypes: float64(4), object(3)
memory usage: 18.3+ KB
```

```
In [ ]: val_count_cols = ['species', 'island', 'sex']
for col in val_count_cols:
    print(penguins[col].value_counts())
    print()

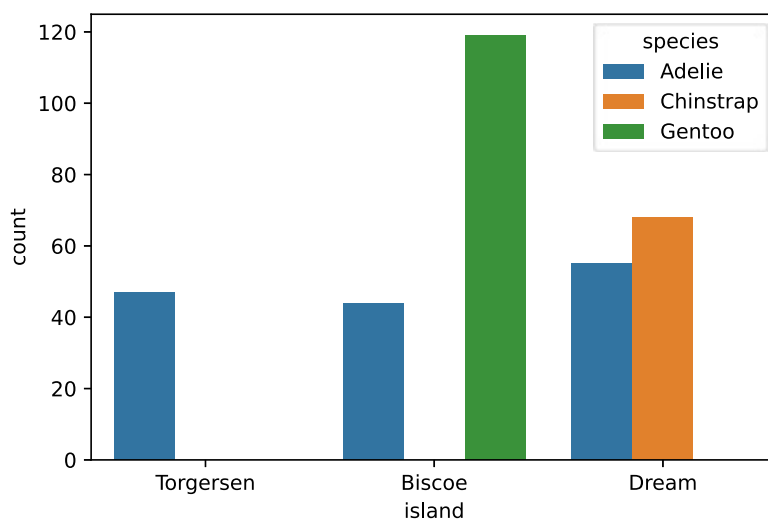
Adelie      146
Gentoo       119
Chinstrap    68
Name: species, dtype: int64

Biscoe      163
Dream        123
Torgersen    47
Name: island, dtype: int64

MALE        168
FEMALE      165
Name: sex, dtype: int64
```

```
In [ ]: sns.countplot(data=penguins, x='island', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='island', ylabel='count'>
```




```
In [ ]: penguins.groupby(['island', 'species']).count()
```

```
Out[ ]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex
island species
```

Biscoe	Adelie	44	44	44	44	44
	Gentoo	119	119	119	119	119
Dream	Adelie	55	55	55	55	55
	Chinstrap	68	68	68	68	68
Torgersen	Adelie	47	47	47	47	47

```
In [ ]: sns.countplot(data=penguins, x='species')
```

```
Out[ ]: <AxesSubplot:xlabel='species', ylabel='count'>
```




```
In [ ]: penguins['species'].value_counts()
```

```
Out[ ]: Adelie      146
Gentoo      119
Chinstrap    68
Name: species, dtype: int64
```

```
In [ ]: sns.countplot(data=penguins, x='sex')
```


```
Out[ ]: <AxesSubplot:xlabel='sex', ylabel='count'>
```



```
In [ ]: sns.countplot(data=penguins, x='sex', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='sex', ylabel='count'>
```





```
In [ ]: penguins.sex.value_counts()
```

```
Out[ ]: MALE      168
        FEMALE    165
        Name: sex, dtype: int64
```

```
In [ ]: penguins.groupby(['sex', 'species'])['sex'].count()
```


```
Out[ ]: sex  species
        FEMALE Adelie      73
          Chinstrap    34
          Gentoo      58
        MALE  Adelie      73
          Chinstrap    34
          Gentoo      61
        Name: sex, dtype: int64
```

```
In [ ]: penguins.corr()
```

```
Out[ ]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g
culmen_length_mm      1.000000      -0.228626      0.653096      0.589451
culmen_depth_mm      -0.228626      1.000000     -0.577792     -0.472016
flipper_length_mm      0.653096     -0.577792      1.000000      0.872979
body_mass_g           0.589451     -0.472016      0.872979      1.000000
```


```
In [ ]: sns.scatterplot(data=penguins, x='body_mass_g', y='flipper_length_mm', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='body_mass_g', ylabel='flipper_length_mm'>
```




```
In [ ]: sns.scatterplot(data=penguins, x='culmen_length_mm', y='culmen_depth_mm', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='culmen_length_mm', ylabel='culmen_depth_mm'>
```



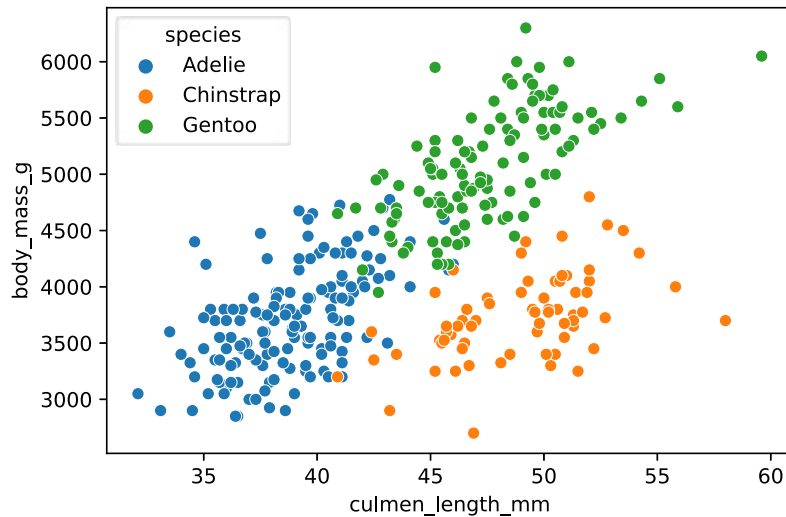
```
In [ ]: sns.scatterplot(data=penguins, x='culmen_length_mm', y='flipper_length_mm', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='culmen_length_mm', ylabel='flipper_length_mm'>
```



```
In [ ]: sns.scatterplot(data=penguins, x='culmen_length_mm', y='body_mass_g', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='culmen_length_mm', ylabel='body_mass_g'>
```



```
In [ ]: px.scatter_3d(penguins, x='culmen_length_mm', y='culmen_depth_mm', z='flipper_length_mm', color='species', title=
```

```
In [ ]: px.scatter_3d(penguins, x='culmen_length_mm', y='culmen_depth_mm', z='body_mass_g', color='species', title='Pengu
```

```
In [ ]: px.scatter_3d(penguins, x='culmen_depth_mm', y='flipper_length_mm', z='body_mass_g', color='species', title='Peng
```

```
In [ ]: px.scatter_3d(penguins, x='culmen_length_mm', y='flipper_length_mm', z='body_mass_g', color='species', title='Per
```



```
In [ ]: penguins = penguins[['island', 'species', 'culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g', 'sex']]
penguins.head()
```



Out[]:

	island	species	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	sex
0	Torgersen	Adelie	39.1	18.7	181.0	3750.0	MALE
1	Torgersen	Adelie	39.5	17.4	186.0	3800.0	FEMALE
2	Torgersen	Adelie	40.3	18.0	195.0	3250.0	FEMALE
3	Torgersen	Adelie	36.7	19.3	193.0	3450.0	FEMALE
4	Torgersen	Adelie	39.3	20.6	190.0	3650.0	MALE

```
In [ ]: plot_cols = ['culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g']

for col in plot_cols:
    sns.displot(data=penguins, x=col, col='species', kind='kde')
```





```
In [ ]: penguins.groupby('species').agg({'culmen_length_mm': ['count', 'min', 'max', 'mean', 'median'], 'culmen_depth_mm':
```

```
Out[ ]:
```

species	culmen_length_mm					culmen_depth_mm					flipper_length_mm				
	count	min	max	mean	median	count	min	max	mean	median	count	min	max	mean	median
Adelie	146	32.1	46.0	38.823973	38.85	146	15.5	21.5	18.347260	18.40	146	172.0	210.0	190.102740	190.0
Chinstrap	68	40.9	58.0	48.833824	49.55	68	16.4	20.8	18.420588	18.45	68	178.0	212.0	195.823529	196.0
Gentoo	119	40.9	59.6	47.568067	47.40	119	13.1	17.3	14.996639	15.00	119	203.0	231.0	217.235294	216.0

```
In [ ]: X = penguins.drop(columns='sex')
y = penguins['sex']
```

```
print(X.shape, y.shape)
```

```
(333, 6) (333,)
```

```
In [ ]: print(y[0])
        print(y[120])

        y = LabelEncoder().fit_transform(y)

        print(y[0])
        print(y[120])
```

```
MALE
FEMALE
1
0
```

```
In [ ]: X = pd.get_dummies(X)
        X.head()
```

```
Out[ ]:   culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  island_Biscoe  island_Dream  island_Torgersen  species_Adelie  sp
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	island_Biscoe	island_Dream	island_Torgersen	species_Adelie	sp
0	39.1	18.7	181.0	3750.0	0	0	1	1	
1	39.5	17.4	186.0	3800.0	0	0	1	1	
2	40.3	18.0	195.0	3250.0	0	0	1	1	
3	36.7	19.3	193.0	3450.0	0	0	1	1	
4	39.3	20.6	190.0	3650.0	0	0	1	1	

Random Forest Classifier -> predict sex

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=7)
        print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)

        (266, 10) (266,) (67, 10) (67,)
```

```
In [ ]: model = RandomForestClassifier(n_estimators=100, criterion='gini', max_depth=6, random_state=7)
        model.fit(X_train, y_train)
```

```
Out[ ]: RandomForestClassifier(max_depth=6, random_state=7)
```

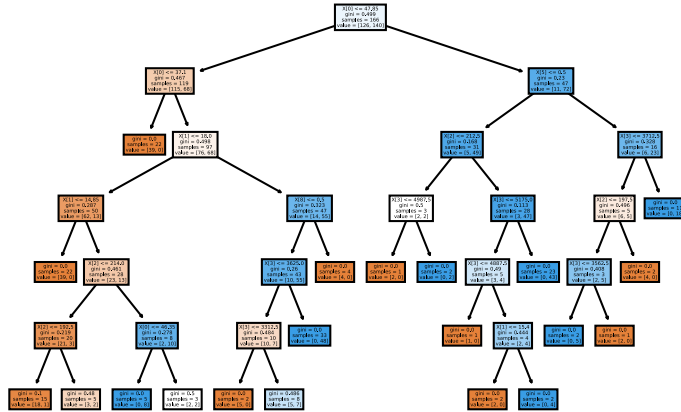
```
In [ ]: estimator = model.estimators_[5]

        plot_tree(estimator, filled = True)
```

```
Out[ ]: [Text(173.6, 201.90857142857143, 'X[0] <= 47.85\ngini = 0.499\nsamples = 166\nvalue = [126, 140]'),
        Text(80.60000000000001, 170.84571428571428, 'X[0] <= 37.1\ngini = 0.467\nsamples = 119\nvalue = [115, 68]'),
        Text(68.2, 139.78285714285715, 'gini = 0.0\nsamples = 22\nvalue = [39, 0]'),
        Text(93.0, 139.78285714285715, 'X[1] <= 18.0\ngini = 0.498\nsamples = 97\nvalue = [76, 68]'),
        Text(37.2, 108.72, 'X[1] <= 14.85\ngini = 0.287\nsamples = 50\nvalue = [62, 13]'),
        Text(24.8, 77.65714285714284, 'gini = 0.0\nsamples = 22\nvalue = [39, 0]'),
        Text(49.6, 77.65714285714284, 'X[2] <= 214.0\ngini = 0.461\nsamples = 28\nvalue = [23, 13]'),
        Text(24.8, 46.59428571428572, 'X[2] <= 192.5\ngini = 0.219\nsamples = 20\nvalue = [21, 3]'),
        Text(12.4, 15.531428571428563, 'gini = 0.1\nsamples = 15\nvalue = [18, 1]'),
        Text(37.2, 15.531428571428563, 'gini = 0.48\nsamples = 5\nvalue = [3, 2]'),
        Text(74.4, 46.59428571428572, 'X[0] <= 46.35\ngini = 0.278\nsamples = 8\nvalue = [2, 10]'),
        Text(62.0, 15.531428571428563, 'gini = 0.0\nsamples = 5\nvalue = [0, 8]'),
        Text(86.8, 15.531428571428563, 'gini = 0.5\nsamples = 3\nvalue = [2, 2]'),
        Text(148.8, 108.72, 'X[8] <= 0.5\ngini = 0.323\nsamples = 47\nvalue = [14, 55]'),
        Text(136.4, 77.65714285714284, 'X[3] <= 3625.0\ngini = 0.26\nsamples = 43\nvalue = [10, 55]'),
        Text(124.0, 46.59428571428572, 'X[3] <= 3312.5\ngini = 0.484\nsamples = 10\nvalue = [10, 7]'),
        Text(111.60000000000001, 15.531428571428563, 'gini = 0.0\nsamples = 2\nvalue = [5, 0]'),
        Text(136.4, 15.531428571428563, 'gini = 0.0\nsamples = 8\nvalue = [5, 7]'),
        Text(148.8, 46.59428571428572, 'gini = 0.0\nsamples = 33\nvalue = [0, 48]'),
        Text(161.20000000000002, 77.65714285714284, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
        Text(266.6, 170.84571428571428, 'X[5] <= 0.5\ngini = 0.23\nsamples = 47\nvalue = [11, 72]'),
        Text(223.20000000000002, 139.78285714285715, 'X[2] <= 212.5\ngini = 0.168\nsamples = 31\nvalue = [5, 49]'),
        Text(198.4, 108.72, 'X[3] <= 4987.5\ngini = 0.5\nsamples = 3\nvalue = [2, 2]'),
        Text(186.0, 77.65714285714284, 'gini = 0.0\nsamples = 1\nvalue = [2, 0]'),
        Text(210.8, 77.65714285714284, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
        Text(248.0, 108.72, 'X[3] <= 5175.0\ngini = 0.113\nsamples = 28\nvalue = [3, 47]'),
```



```
Text(235.6, 77.65714285714284, 'X[3] <= 4887.5\ngini = 0.49\nsamples = 5\nvalue = [3, 4]'),
Text(223.20000000000002, 46.59428571428572, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(248.0, 46.59428571428572, 'X[1] <= 15.4\ngini = 0.444\nsamples = 4\nvalue = [2, 4]'),
Text(235.6, 15.531428571428563, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(260.40000000000003, 15.531428571428563, 'gini = 0.0\nsamples = 2\nvalue = [0, 4]'),
Text(260.40000000000003, 77.65714285714284, 'gini = 0.0\nsamples = 23\nvalue = [0, 43]'),
Text(310.0, 139.78285714285715, 'X[3] <= 3712.5\ngini = 0.328\nsamples = 16\nvalue = [6, 23]'),
Text(297.6, 108.72, 'X[2] <= 197.5\ngini = 0.496\nsamples = 5\nvalue = [6, 5]'),
Text(285.2, 77.65714285714284, 'X[3] <= 3562.5\ngini = 0.408\nsamples = 3\nvalue = [2, 5]'),
Text(272.8, 46.59428571428572, 'gini = 0.0\nsamples = 2\nvalue = [0, 5]'),
Text(297.6, 46.59428571428572, 'gini = 0.0\nsamples = 1\nvalue = [2, 0]'),
Text(310.0, 77.65714285714284, 'gini = 0.0\nsamples = 2\nvalue = [4, 0]'),
Text(322.40000000000003, 108.72, 'gini = 0.0\nsamples = 11\nvalue = [0, 18]')]
```




```
In [ ]: y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)
```

```
In [ ]: print(classification_report(y_train, y_pred_train))
```

	precision	recall	f1-score	support
0	0.98	0.99	0.99	127
1	0.99	0.99	0.99	139
accuracy			0.99	266
macro avg	0.99	0.99	0.99	266
weighted avg	0.99	0.99	0.99	266

```
In [ ]: plot_confusion_matrix(model, X_train, y_train, display_labels=['FEMALE', 'MALE'])
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1fbf2756520>
```



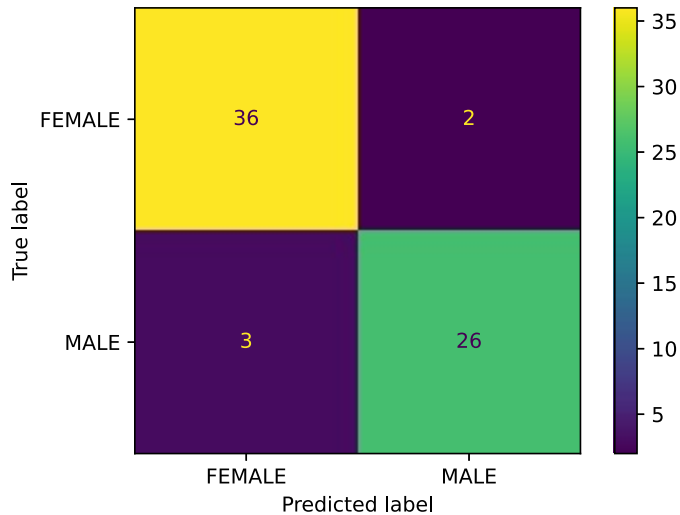
```
In [ ]: print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.92	0.95	0.94	38
1	0.93	0.90	0.91	29
accuracy			0.93	67
macro avg	0.93	0.92	0.92	67
weighted avg	0.93	0.93	0.93	67

```
In [ ]: plot_confusion_matrix(model, X_test, y_test, display_labels=['FEMALE', 'MALE'])
```

```
Out[ ]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1fbf14ede50>
```



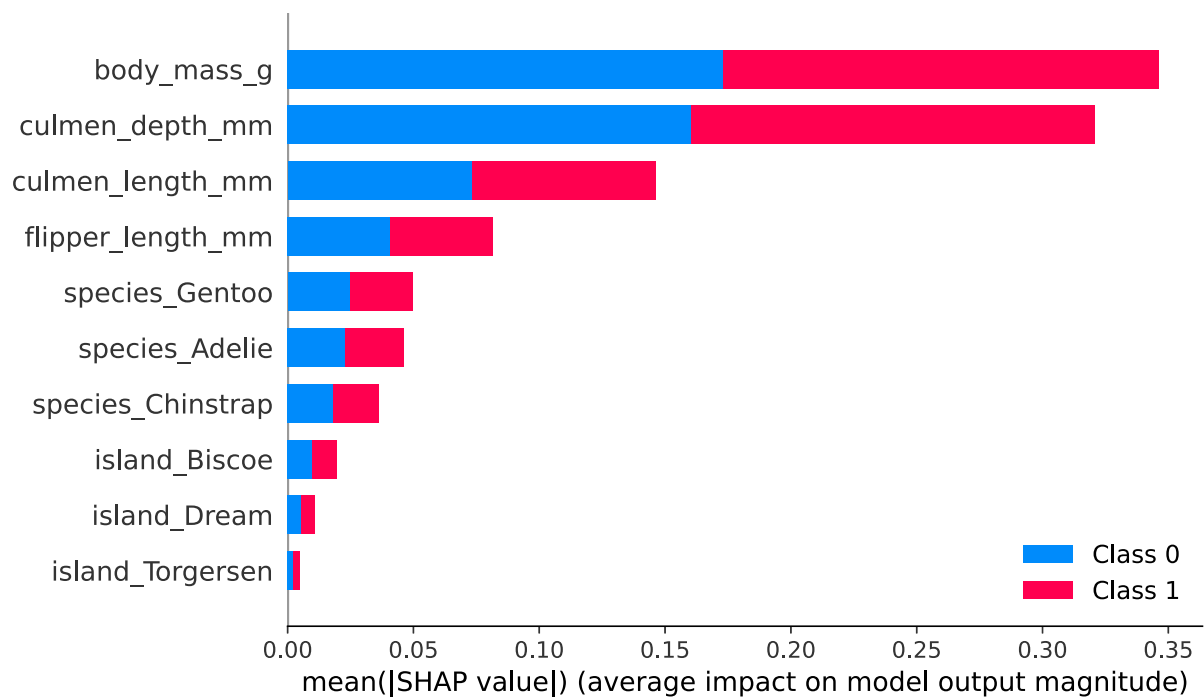
```
In [ ]: y_pred_proba = model.predict_proba(X_test)
y_pred_proba
```

```
Out[ ]: array([[9.89259329e-01, 1.07406711e-02],
 [7.58275514e-02, 9.24172449e-01],
 [9.16221287e-02, 9.08377871e-01],
 [2.82487259e-02, 9.71751274e-01],
 [2.32558140e-04, 9.99767442e-01],
 [9.32253144e-01, 6.77468562e-02],
 [9.73028779e-01, 2.69712211e-02],
 [1.99836185e-02, 9.80016381e-01],
 [9.65358230e-01, 3.46417700e-02],
 [9.31472444e-01, 6.85275562e-02],
 [9.92937595e-01, 7.06240495e-03],
 [2.32558140e-04, 9.99767442e-01],
 [3.27363366e-01, 6.72636634e-01],
 [9.74020928e-01, 2.59790716e-02],
 [5.63459880e-02, 9.43654012e-01],
 [9.91662801e-01, 8.33719887e-03],
 [1.04634463e-01, 8.95365537e-01],
 [7.55851968e-01, 2.44148032e-01],
 [2.32558140e-04, 9.99767442e-01],
 [3.71320084e-02, 9.62867992e-01],
 [2.41802044e-01, 7.58197956e-01],
 [9.90936874e-01, 9.06312645e-03],
 [7.42169531e-01, 2.57830469e-01],
 [9.74910912e-01, 2.50890877e-02],
 [6.90924400e-01, 3.09075600e-01],
 [3.31674499e-01, 6.68325501e-01],
 [9.81214906e-01, 1.87850940e-02],
 [2.32558140e-04, 9.99767442e-01],
 [8.68116075e-01, 1.31883925e-01],
 [9.31979856e-01, 6.80201439e-02],
 [4.69294647e-01, 5.30705353e-01],
 [7.99055667e-01, 2.00944333e-01],
 [9.52142851e-01, 4.78571490e-02],
 [9.30987469e-01, 6.90125312e-02],
 [5.94129565e-01, 4.05870435e-01],
 [3.95260586e-01, 6.04739414e-01],
 [9.89500095e-01, 1.04999050e-02],
 [4.05742284e-02, 9.59425772e-01],
 [5.03897251e-01, 4.96102749e-01],
 [4.01636754e-01, 5.98363246e-01],
 [9.47727368e-01, 5.22726323e-02],
```


```
[9.73248110e-01, 2.67518899e-02],
[8.49187832e-01, 1.50812168e-01],
[7.92965954e-01, 2.07034046e-01],
[1.70951253e-02, 9.82904875e-01],
[9.45799971e-01, 5.42000285e-02],
[9.80412801e-01, 1.95871989e-02],
[9.55511968e-01, 4.44880318e-02],
[7.41874325e-01, 2.58125675e-01],
[9.86893257e-01, 1.31067426e-02],
[3.53769841e-02, 9.64623016e-01],
[2.75530303e-01, 7.24469697e-01],
[3.95363170e-01, 6.04636830e-01],
[6.34111888e-01, 3.65888112e-01],
[9.91866166e-01, 8.13383352e-03],
[9.84092662e-01, 1.59073378e-02],
[7.72590208e-01, 2.27409792e-01],
[8.14668978e-01, 1.85331022e-01],
[5.24670855e-02, 9.47532914e-01],
[3.76105305e-01, 6.23894695e-01],
[2.32558140e-04, 9.99767442e-01],
[1.87646663e-01, 8.12353337e-01],
[3.50015813e-01, 6.49984187e-01],
[9.57614689e-01, 4.23853107e-02],
[1.22232046e-01, 8.77767954e-01],
[9.78012878e-01, 2.19871224e-02],
[5.34451167e-01, 4.65548833e-01]]])
```


```
In [ ]: explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_test)
```

```
In [ ]: shap.summary_plot(shap_values, X_test)
```




```
In [ ]: shap.summary_plot(shap_values[0], X_test)
```







```
In [ ]: shap.summary_plot(shap_values[1], X_test)
```




```
In [ ]: shap.decision_plot(explainer.expected_value[0], shap_values[0], X_test.iloc[0, :])
```




```
In [ ]: shap.decision_plot(explainer.expected_value[1], shap_values[1], X_test.iloc[1, :])
```







```
In [ ]: shap.dependence_plot('body_mass_g', shap_values[0], X_test)
```



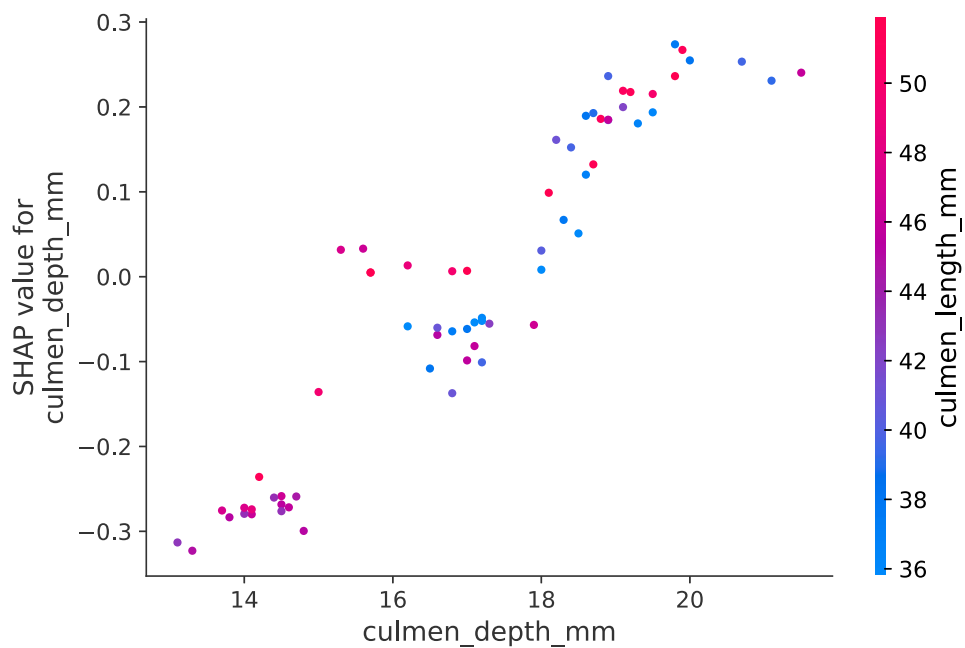
```
In [ ]: shap.dependence_plot('body_mass_g', shap_values[1], X_test)
```




```
In [ ]: shap.dependence_plot('culmen_depth_mm', shap_values[0], X_test)
```




```
In [ ]: shap.dependence_plot('culmen_depth_mm', shap_values[1], X_test)
```




```
In [ ]: shap.dependence_plot('culmen_length_mm', shap_values[0], X_test)
```

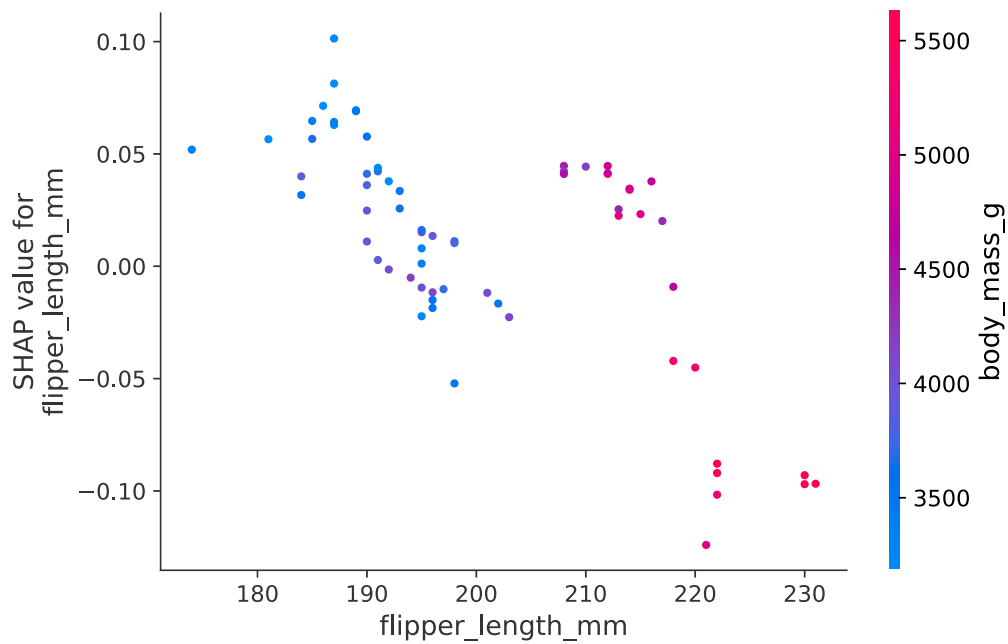





```
In [ ]: shap.dependence_plot('culmen_length_mm', shap_values[1], X_test)
```




```
In [ ]: shap.dependence_plot('flipper_length_mm', shap_values[0], X_test)
```




```
In [ ]: shap.dependence_plot('flipper_length_mm', shap_values[1], X_test)
```






```
In [ ]: shap.dependence_plot('species_Gentoo', shap_values[0], X_test)
```




```
In [ ]: shap.dependence_plot('species_Gentoo', shap_values[1], X_test)
```




```
In [ ]: shap.dependence_plot('species_Adelie', shap_values[0], X_test)
```






```
In [ ]: shap.dependence_plot('species_Adelie', shap_values[1], X_test)
```



```
In [ ]: shap.dependence_plot('species_Chinstrap', shap_values[0], X_test)
```



```
In [ ]: shap.initjs()
shap.force_plot(explainer.expected_value[0], shap_values[0], X_test.iloc[:, :])
```

Out[]:






```
In [ ]: shap.force_plot(explainer.expected_value[1], shap_values[1], X_test.iloc[:, :])
```

Out[]:

```
In [ ]: shap_values_test0 = explainer.shap_values(X_test.iloc[0])
shap.force_plot(explainer.expected_value[0], shap_values_test0[0], X_test.iloc[0])
```

Out[]:


```
In [ ]: shap.decision_plot(explainer.expected_value[0], shap_values[0][0], X_test.iloc[0, :])
```



```
In [ ]: shap.force_plot(explainer.expected_value[1], shap_values_test0[1], X_test.iloc[0])
```

Out[]:


```
In [ ]: shap.decision_plot(explainer.expected_value[1], shap_values[1][0], X_test.iloc[0, :])
```



```
In [ ]: shap_values_test3 = explainer.shap_values(X_test.iloc[3])
shap.force_plot(explainer.expected_value[0], shap_values_test3[0], X_test.iloc[3])
```

Out[]:


```
In [ ]: shap.decision_plot(explainer.expected_value[0], shap_values[0][3], X_test.iloc[3, :])
```



```
In [ ]: shap.force_plot(explainer.expected_value[1], shap_values_test3[1], X_test.iloc[3])
```

Out[]:


```
In [ ]: shap.decision_plot(explainer.expected_value[1], shap_values[1][3], X_test.iloc[3, :])
```




```
In [ ]: shap_values_test12 = explainer.shap_values(X_test.iloc[12])
shap.force_plot(explainer.expected_value[0], shap_values_test12[0], X_test.iloc[12])
```

Out[]:

```
In [ ]: shap.decision_plot(explainer.expected_value[0], shap_values[0][12], X_test.iloc[12, :])
```




island_Biscoe



```
In [ ]: shap.force_plot(explainer.expected_value[1], shap_values_test12[1], X_test.iloc[12])
```

Out[]:

```
In [ ]: shap.decision_plot(explainer.expected_value[1], shap_values[1][12], X_test.iloc[12, :])
```



```
In [ ]:
```