

# Machine Learning - Explain Models with SHAP

Classification

<https://www.kaggle.com/parulpandey/palmer-archipelago-antarctica-penguin-data>

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from sklearn.preprocessing import StandardScaler
from sklearn.tree import plot_tree
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import shap
```

```
In [ ]: penguins = pd.read_csv('penguins_size.csv')
penguins.head()
```

```
Out[ ]:   species  island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex
0   Adelie  Torgersen             39.1             18.7             181.0         3750.0  MALE
1   Adelie  Torgersen             39.5             17.4             186.0         3800.0  FEMALE
2   Adelie  Torgersen             40.3             18.0             195.0         3250.0  FEMALE
3   Adelie  Torgersen             NaN             NaN             NaN             NaN     NaN
4   Adelie  Torgersen             36.7             19.3             193.0         3450.0  FEMALE
```

```
In [ ]: penguins.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   species                344 non-null   object  
 1   island                 344 non-null   object  
 2   culmen_length_mm       342 non-null   float64  
 3   culmen_depth_mm        342 non-null   float64  
 4   flipper_length_mm      342 non-null   float64  
 5   body_mass_g            342 non-null   float64  
 6   sex                    334 non-null   object  
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

```
In [ ]: penguins.describe(include='all')
```

```
Out[ ]:   species  island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex
count      344     344           342.000000         342.000000         342.000000         342.000000     334
unique        3        3              NaN              NaN              NaN              NaN         3
top    Adelie  Biscoe              NaN              NaN              NaN              NaN     MALE
freq      152     168              NaN              NaN              NaN              NaN     168
mean        NaN     NaN           43.921930          17.151170          200.915205        4201.754386     NaN
std         NaN     NaN           5.459584           1.974793           14.061714          801.954536     NaN
min         NaN     NaN           32.100000          13.100000          172.000000        2700.000000     NaN
25%         NaN     NaN           39.225000          15.600000          190.000000        3550.000000     NaN
50%         NaN     NaN           44.450000          17.300000          197.000000        4050.000000     NaN
75%         NaN     NaN           48.500000          18.700000          213.000000        4750.000000     NaN
max         NaN     NaN           59.600000          21.500000          231.000000        6300.000000     NaN
```

```
In [ ]: penguins.isna().sum()
```

```
Out[ ]: species      0
island      0
culmen_length_mm    2
```

```
culmen_depth_mm      2
flipper_length_mm    2
body_mass_g          2
sex                  10
dtype: int64
```

```
In [ ]: penguins = penguins.dropna()
penguins = penguins[penguins['sex'] != '.']
penguins = penguins.reset_index(drop=True)
penguins.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 333 entries, 0 to 332
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   species         333 non-null   object
1   island          333 non-null   object
2   culmen_length_mm 333 non-null   float64
3   culmen_depth_mm  333 non-null   float64
4   flipper_length_mm 333 non-null   float64
5   body_mass_g      333 non-null   float64
6   sex             333 non-null   object
dtypes: float64(4), object(3)
memory usage: 18.3+ KB
```

## Exploratory Data Analysis

```
In [ ]: val_count_cols = ['species', 'island', 'sex']
for col in val_count_cols:
    print(penguins[col].value_counts())
    print()
```

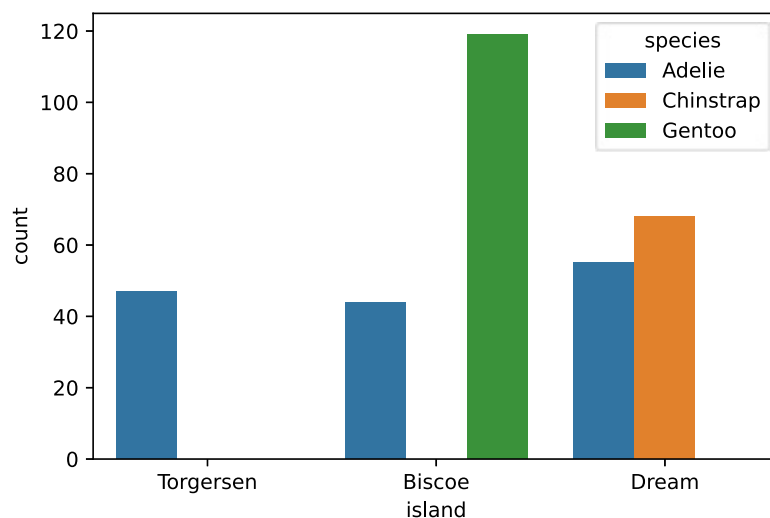
```
Adelie      146
Gentoo      119
Chinstrap   68
Name: species, dtype: int64
```

```
Biscoe      163
Dream       123
Torgersen    47
Name: island, dtype: int64
```

```
MALE      168
FEMALE    165
Name: sex, dtype: int64
```

```
In [ ]: sns.countplot(data=penguins, x='island', hue='species')
```


```
Out[ ]: <AxesSubplot:xlabel='island', ylabel='count'>
```



```
In [ ]: sns.countplot(data=penguins, x='species')
```


```
In [ ]: sns.countplot(data=penguins, x='species')
```

```
Out[ ]: <AxesSubplot:xlabel='species', ylabel='count'>
```



```
In [ ]: sns.countplot(data=penguins, x='sex')
```

```
Out[ ]: <AxesSubplot:xlabel='sex', ylabel='count'>
```



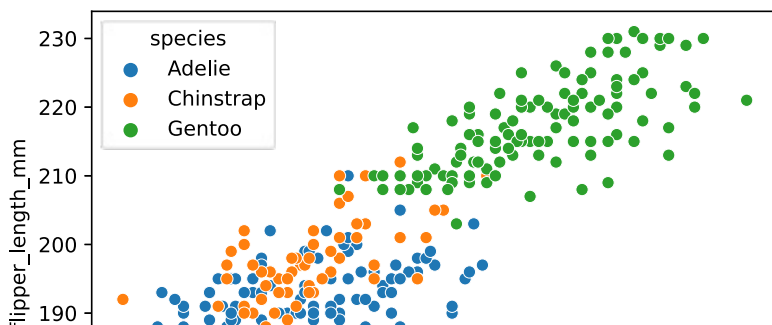
```
In [ ]: penguins.corr()
```


```
Out[ ]:
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
culmen_length_mm	1.000000	-0.228626	0.653096	0.589451
culmen_depth_mm	-0.228626	1.000000	-0.577792	-0.472016
flipper_length_mm	0.653096	-0.577792	1.000000	0.872979
body_mass_g	0.589451	-0.472016	0.872979	1.000000

```
In [ ]: sns.scatterplot(data=penguins, x='body_mass_g', y='flipper_length_mm', hue='species')
```


```
Out[ ]: <AxesSubplot:xlabel='body_mass_g', ylabel='flipper_length_mm'>
```






```
In [ ]: sns.scatterplot(data=penguins, x='culmen_length_mm', y='culmen_depth_mm', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='culmen_length_mm', ylabel='culmen_depth_mm'>
```




```
In [ ]: sns.scatterplot(data=penguins, x='culmen_length_mm', y='flipper_length_mm', hue='species')
```


```
Out[ ]: <AxesSubplot:xlabel='culmen_length_mm', ylabel='flipper_length_mm'>
```



```
In [ ]: sns.scatterplot(data=penguins, x='culmen_length_mm', y='body_mass_g', hue='species')
```

```
Out[ ]: <AxesSubplot:xlabel='culmen_length_mm', ylabel='body_mass_g'>
```





```
In [ ]: px.scatter_3d(penguins, x='culmen_length_mm', y='culmen_depth_mm', z='flipper_length_mm', color='species', title='Penguins 3D Scatter Plot')
```

```
In [ ]: px.scatter_3d(penguins, x='culmen_length_mm', y='culmen_depth_mm', z='body_mass_g', color='species', title='Penguins 3D Scatter Plot')
```

```
In [ ]: px.scatter_3d(penguins, x='culmen_depth_mm', y='flipper_length_mm', z='body_mass_g', color='species', title='Penguins 3D Scatter Plot')
```

```
In [ ]: px.scatter_3d(penguins, x='culmen_length_mm', y='flipper_length_mm', z='body_mass_g', color='species', title='Penguins 3D Scatter Plot')
```


```
In [ ]: penguins = penguins[['island', 'species', 'culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g', 'sex']]
penguins.head()
```




	island	species	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	sex
0	Torgersen	Adelie	39.1	18.7	181.0	3750.0	MALE
1	Torgersen	Adelie	39.5	17.4	186.0	3800.0	FEMALE
2	Torgersen	Adelie	40.3	18.0	195.0	3250.0	FEMALE
3	Torgersen	Adelie	36.7	19.3	193.0	3450.0	FEMALE
4	Torgersen	Adelie	39.3	20.6	190.0	3650.0	MALE

```
In [ ]: # plot distributions of numerical columns by species

plot_cols = ['culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g']

for col in plot_cols:
    sns.displot(data=penguins, x=col, col='species', kind='kde')
```





```
In [ ]: # statistics for each species
```

```
penguins.groupby('species').agg({'culmen_length_mm': ['count', 'min', 'max', 'mean', 'median'], 'culmen_depth_mm': ['count', 'min', 'max', 'mean', 'median']})
```

	culmen_length_mm					culmen_depth_mm					flipper_length_mm						
	count	min	max	mean	median	count	min	max	mean	median	count	min	max	mean	median	count	min
species																	
Adelie	146	32.1	46.0	38.823973	38.85	146	15.5	21.5	18.347260	18.40	146	172.0	210.0	190.102740	190.0	146	2850.0
Chinstrap	68	40.9	58.0	48.833824	49.55	68	16.4	20.8	18.420588	18.45	68	178.0	212.0	195.823529	196.0	68	2700.0
Gentoo	119	40.9	59.6	47.568067	47.40	119	13.1	17.3	14.996639	15.00	119	203.0	231.0	217.235294	216.0	119	3950.0

```
In [ ]: # one hot encoding for categoricals
```

```
penguins = pd.get_dummies(penguins)
penguins.head()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	island_Biscoe	island_Dream	island_Torgersen	species_Adelie	species_Chinstrap	species_Gentoo
0	39.1	18.7	181.0	3750.0	0	0	1	1	0	0
1	39.5	17.4	186.0	3800.0	0	0	1	1	0	0
2	40.3	18.0	195.0	3250.0	0	0	1	1	0	0
3	36.7	19.3	193.0	3450.0	0	0	1	1	0	0
4	39.3	20.6	190.0	3650.0	0	0	1	1	0	0

```
In [ ]: # independent and target variables
```

```
X = penguins.drop(columns='body_mass_g')
y = penguins['body_mass_g']
```

```
print(X.shape, y.shape)
```

```
(333, 11) (333,)
```

## Random Forest Regressor -> predict body mass

```
In [ ]: # split dataset into train and test
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=7)
print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)
```

```
(266, 11) (266,) (67, 11) (67,)
```

```
In [ ]: # model
```

```
model = RandomForestRegressor(n_estimators=100, criterion='mse', max_depth=20, min_samples_split=30, max_features='sqrt')
model.fit(X_train, y_train)
```

```
Out[ ]: RandomForestRegressor(max_depth=20, max_features=1, min_samples_split=30,
                             random_state=7)
```

```
In [ ]: # plot for tree number 5 of the random forest model
```

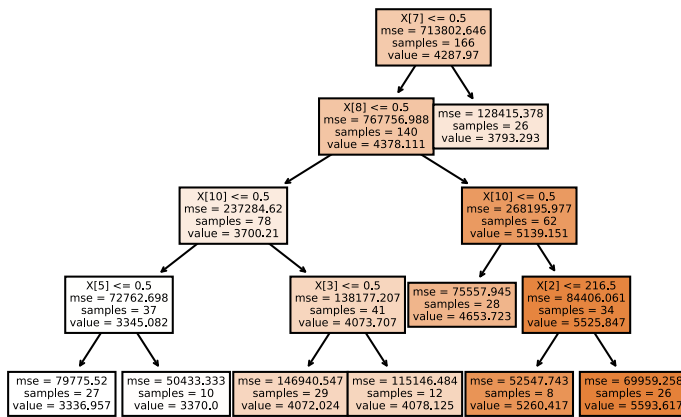
```
from sklearn.tree import plot_tree
```

```
estimator = model.estimators_[5]
```

```
plot_tree(estimator, filled = True)
```

```
Out[ ]: [Text(209.25000000000003, 195.696, 'X[7] <= 0.5\nmse = 713802.646\nsamples = 166\nvalue = 4287.97'),
Text(181.35000000000002, 152.208, 'X[8] <= 0.5\nmse = 767756.988\nsamples = 140\nvalue = 4378.111'),
Text(111.60000000000001, 108.72, 'X[10] <= 0.5\nmse = 237284.62\nsamples = 78\nvalue = 3700.21'),
Text(55.800000000000004, 65.232, 'X[5] <= 0.5\nmse = 72762.698\nsamples = 37\nvalue = 3345.082'),
Text(27.900000000000002, 21.744, 'mse = 79775.52\nsamples = 27\nvalue = 3336.957'),
Text(83.7, 21.744, 'mse = 50433.333\nsamples = 10\nvalue = 3370.0'),
Text(167.4, 65.232, 'X[3] <= 0.5\nmse = 138177.207\nsamples = 41\nvalue = 4073.707'),
Text(139.5, 21.744, 'mse = 146940.547\nsamples = 29\nvalue = 4072.024'),
Text(195.3, 21.744, 'mse = 115146.484\nsamples = 12\nvalue = 4078.125'),
Text(251.10000000000002, 108.72, 'X[10] <= 0.5\nmse = 268195.977\nsamples = 62\nvalue = 5139.151'),
Text(223.20000000000002, 65.232, 'mse = 75557.945\nsamples = 28\nvalue = 4653.723'),
```

```
Text(279.0, 65.232, 'X[2] <= 216.5\nmse = 84406.061\nsamples = 34\nvalue = 5525.847'),
Text(251.10000000000002, 21.744, 'mse = 52547.743\nsamples = 8\nvalue = 5260.417'),
Text(306.90000000000003, 21.744, 'mse = 69959.258\nsamples = 26\nvalue = 5593.617'),
Text(237.15, 152.208, 'mse = 128415.378\nsamples = 26\nvalue = 3793.293')]
```



```
In [ ]: # predictions for train and test sets
```

```
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)
```

```
In [ ]: # performance metrics for train and test sets
```

```
mae_train = mean_absolute_error(y_train, y_pred_train)
mse_train = mean_squared_error(y_train, y_pred_train)
r2_train = r2_score(y_train, y_pred_train)

mae_test = mean_absolute_error(y_test, y_pred_test)
mse_test = mean_squared_error(y_test, y_pred_test)
r2_test = r2_score(y_test, y_pred_test)

print(f'Train: MAE={mae_train} / MSE={mse_train} / R2: {r2_train}')
print(f'Test: MAE={mae_test} / MSE={mse_test} / R2: {r2_test}')
```


```
Train: MAE=227.24073991184062 / MSE=86006.69494581848 / R2: 0.869052917083086
Test: MAE=227.20652137525067 / MSE=78749.41899982275 / R2: 0.8669440122965669
```

## Shap explainers


```
In [ ]: explainer = shap.Explainer(model)
shap_values1 = explainer(X_test)
expected_value = explainer.expected_value

shap_values2 = shap_values = shap.TreeExplainer(model).shap_values(X_test)
```

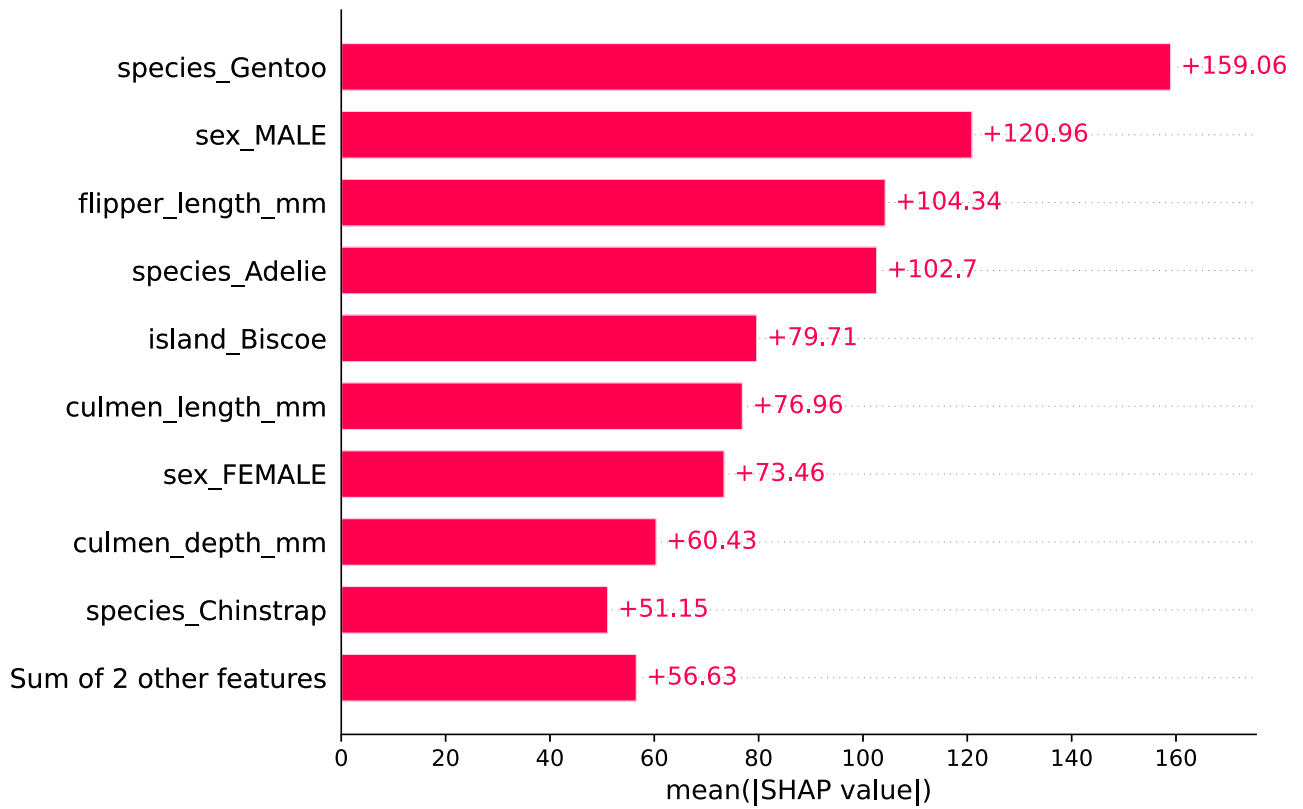
```
In [ ]: shap.plots.beeswarm(shap_values1)
```









```
In [ ]: shap.plots.bar(shap_values1)
```

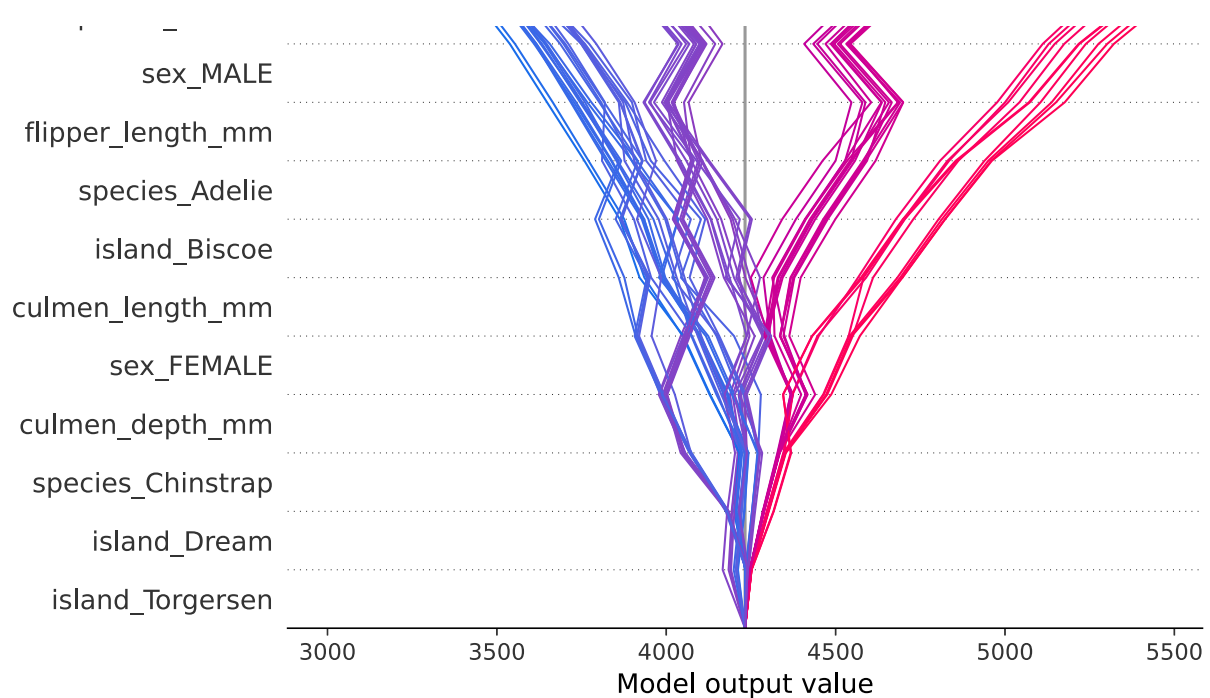


```
In [ ]: shap.plots.heatmap(shap_values1)
```




```
In [ ]: shap.decision_plot(expected_value, shap_values2[:, :], X_test.iloc[:, :])
```






```
In [ ]: shap.dependence_plot('culmen_length_mm', shap_values2, X_test)
```

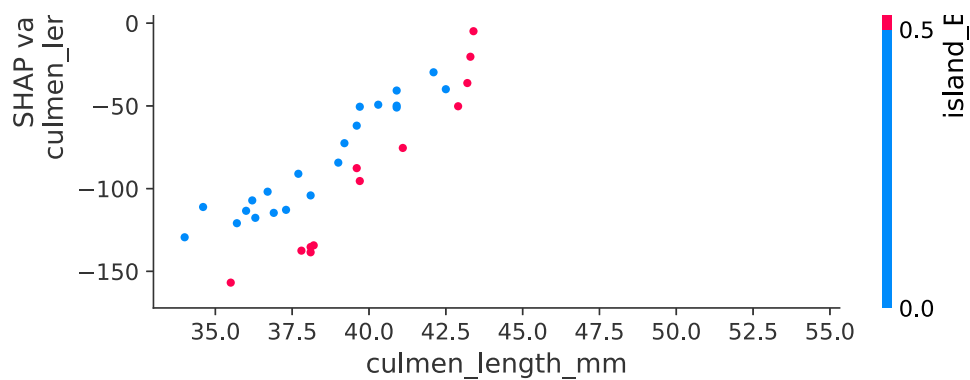
C:\Users\gulia\anaconda3\envs\tpot\lib\site-packages\shap\plots\\_scatter.py:642: MatplotlibDeprecationWarning:  
Passing parameters norm and vmin/vmax simultaneously is deprecated since 3.3 and will become an error two minor releases later. Please pass vmin/vmax directly to the norm when creating it.



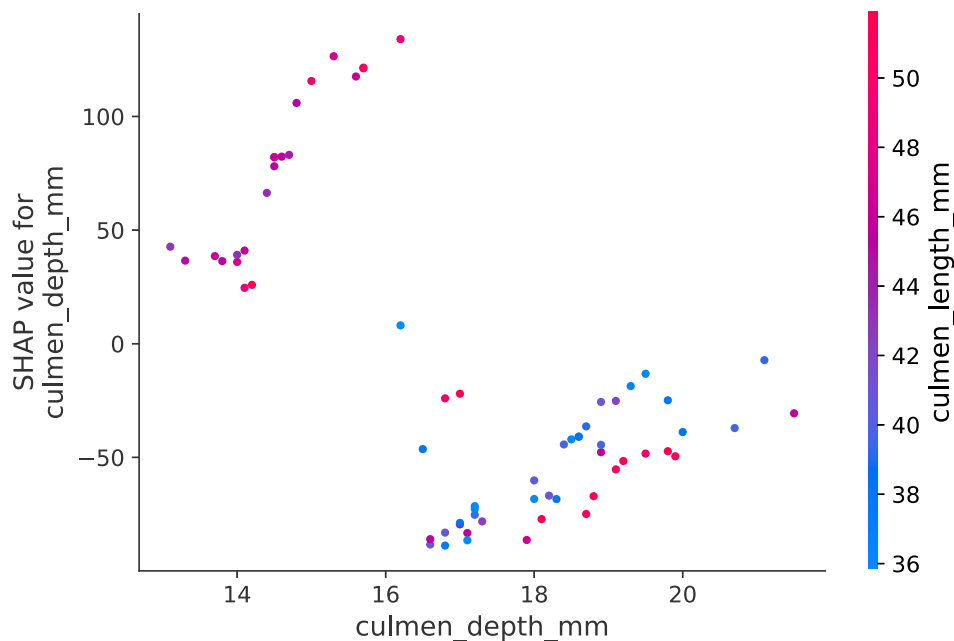
```
In [ ]: shap.dependence_plot('culmen_length_mm', shap_values2, X_test)
```

C:\Users\gulia\anaconda3\envs\tpot\lib\site-packages\shap\plots\\_scatter.py:642: MatplotlibDeprecationWarning:  
Passing parameters norm and vmin/vmax simultaneously is deprecated since 3.3 and will become an error two minor releases later. Please pass vmin/vmax directly to the norm when creating it.






```
In [ ]: shap.dependence_plot('culmen_depth_mm', shap_values2, X_test)
```



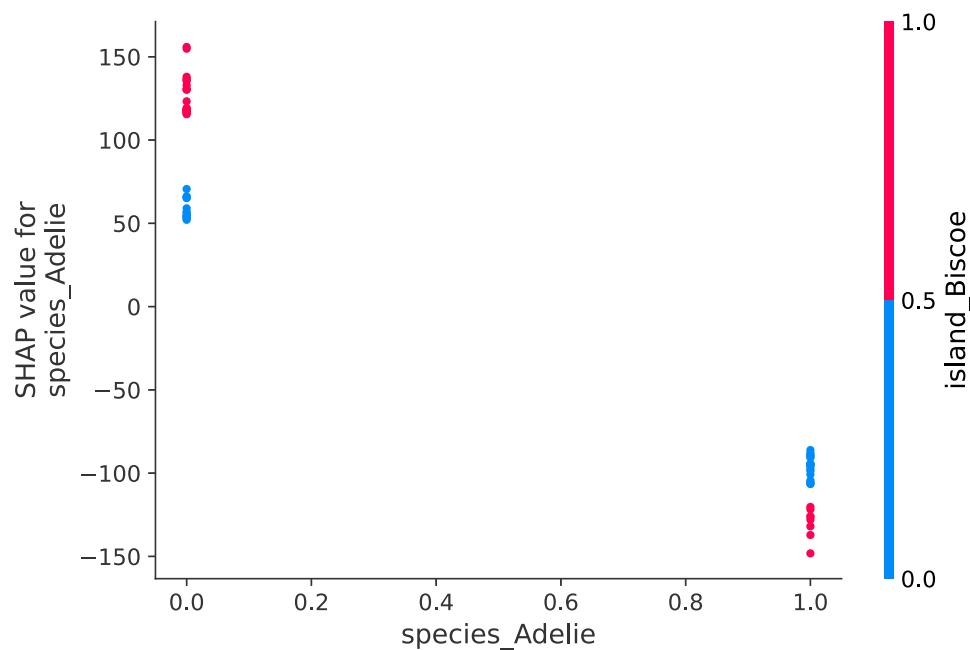
```
In [ ]: shap.dependence_plot('species_Gentoo', shap_values2, X_test)
```



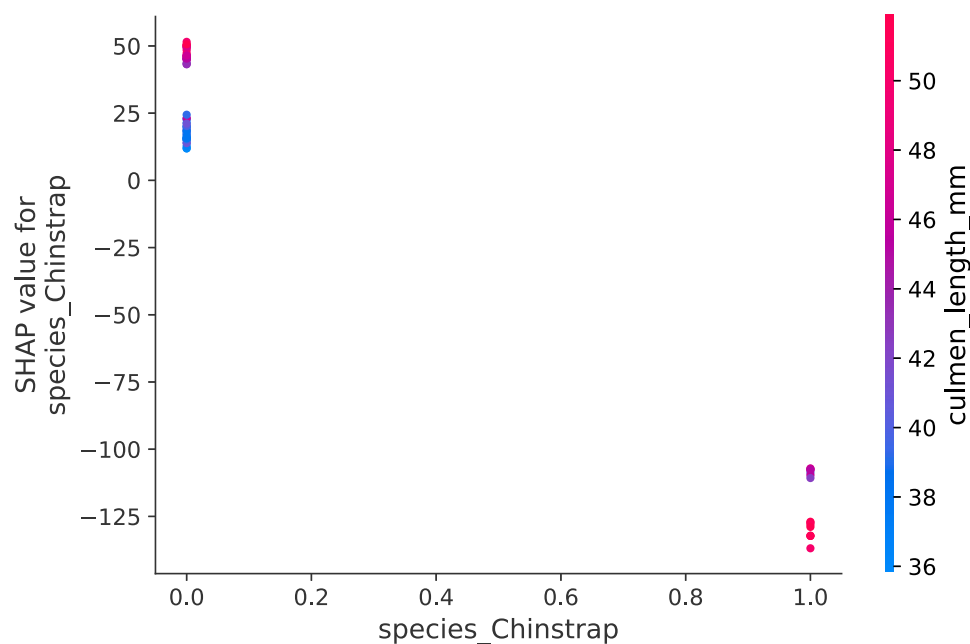
```
In [ ]: shap.dependence_plot('species_Adelie', shap_values2, X_test)
```

C:\Users\gulia\anaconda3\envs\tpot\lib\site-packages\shap\plots\\_scatter.py:642: MatplotlibDeprecationWarning:  
Passing parameters norm and vmin/vmax simultaneously is deprecated since 3.3 and will become an error two minor r

releases later. Please pass vmin/vmax directly to the norm when creating it.




```
In [ ]: shap.dependence_plot('species_Chinstrap', shap_values2, X_test)
```



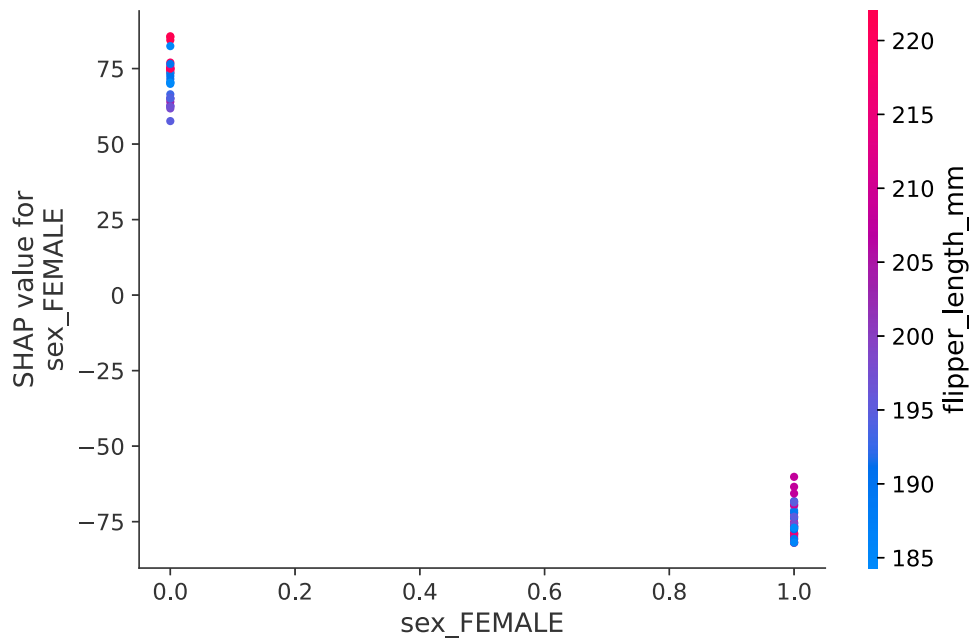
```
In [ ]: shap.dependence_plot('sex_MALE', shap_values2, X_test)
```

C:\Users\gulia\anaconda3\envs\tpot\lib\site-packages\shap\plots\\_scatter.py:642: MatplotlibDeprecationWarning:  
Passing parameters norm and vmin/vmax simultaneously is deprecated since 3.3 and will become an error two minor r  
eleases later. Please pass vmin/vmax directly to the norm when creating it.





```
In [ ]: shap.dependence_plot('sex_FEMALE', shap_values2, X_test)
```



```
In [ ]: shap.initjs()

def p(j):
    return(shap.force_plot(explainer.expected_value, shap_values2[j,:], X_test.iloc[j,:]))
```



```
In [ ]: shap.force_plot(expected_value, shap_values2, X_test.iloc[:, :])
```

```
Out[ ]: 
```

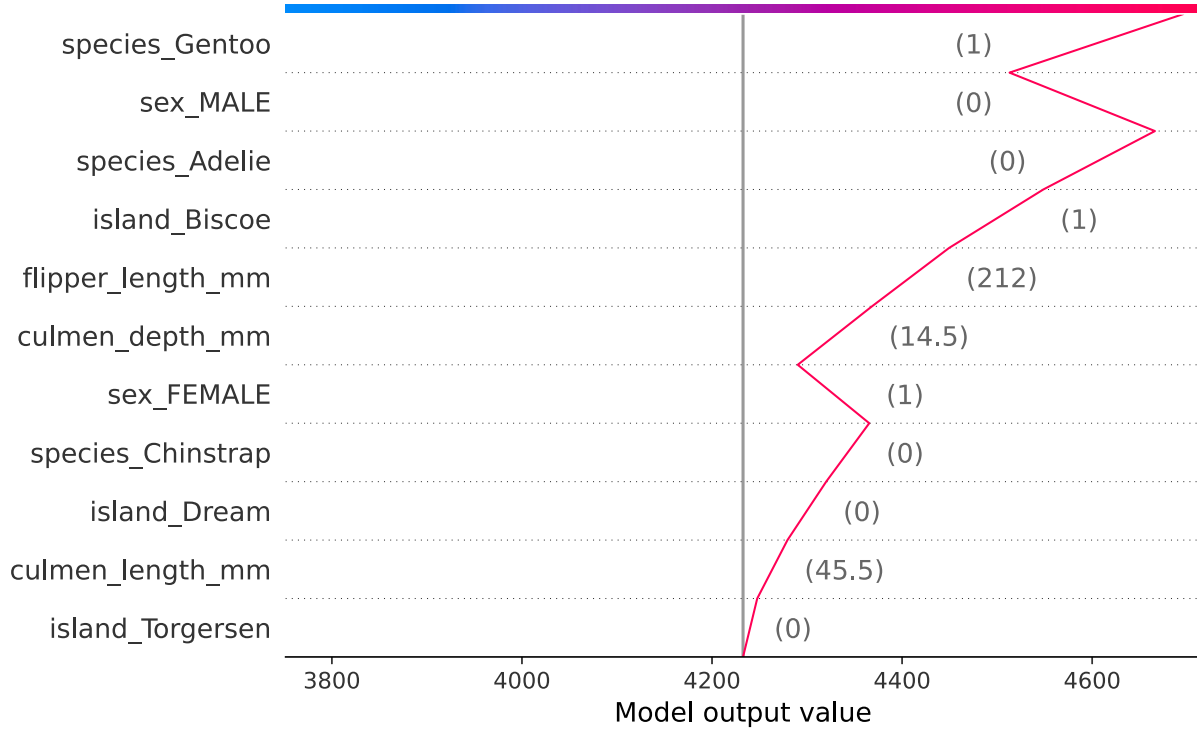


```
In [ ]: p(0)
```

```
Out[ ]:
```

```
In [ ]: shap.decision_plot(expected_value, shap_values2[0, :], X_test.iloc[0, :])
```

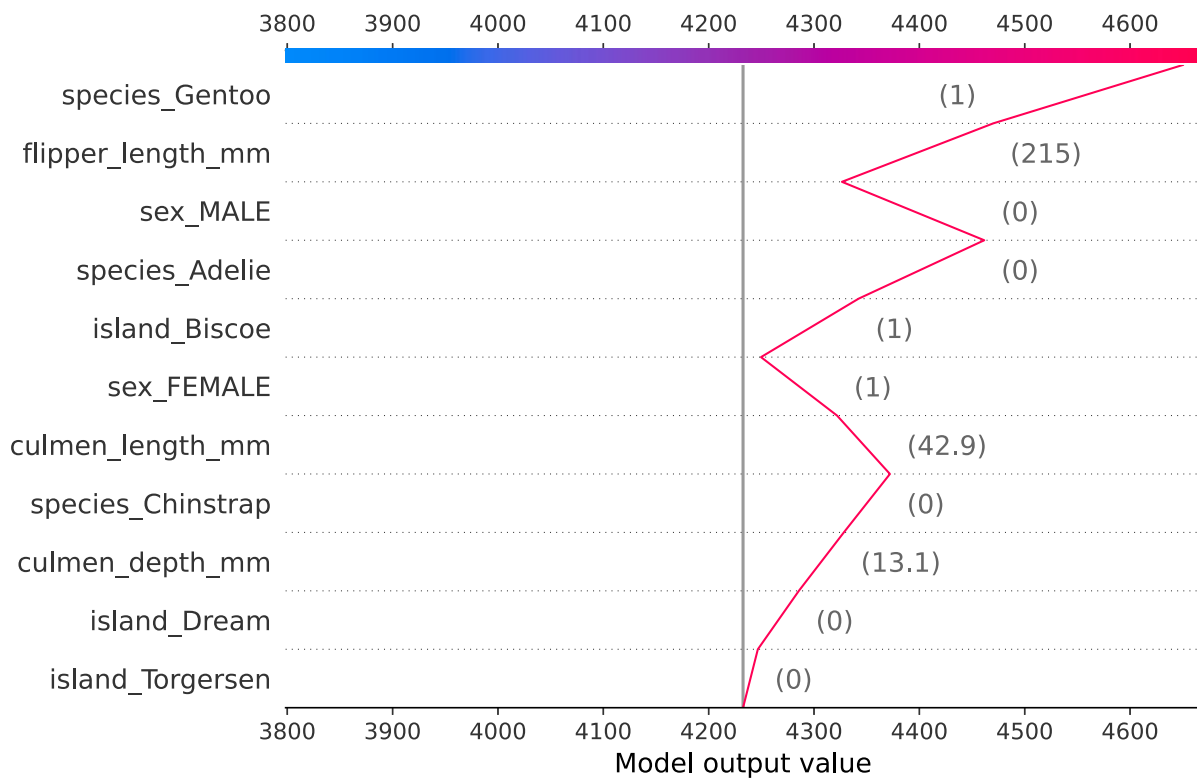




In [ ]: p(33)

Out[ ]:

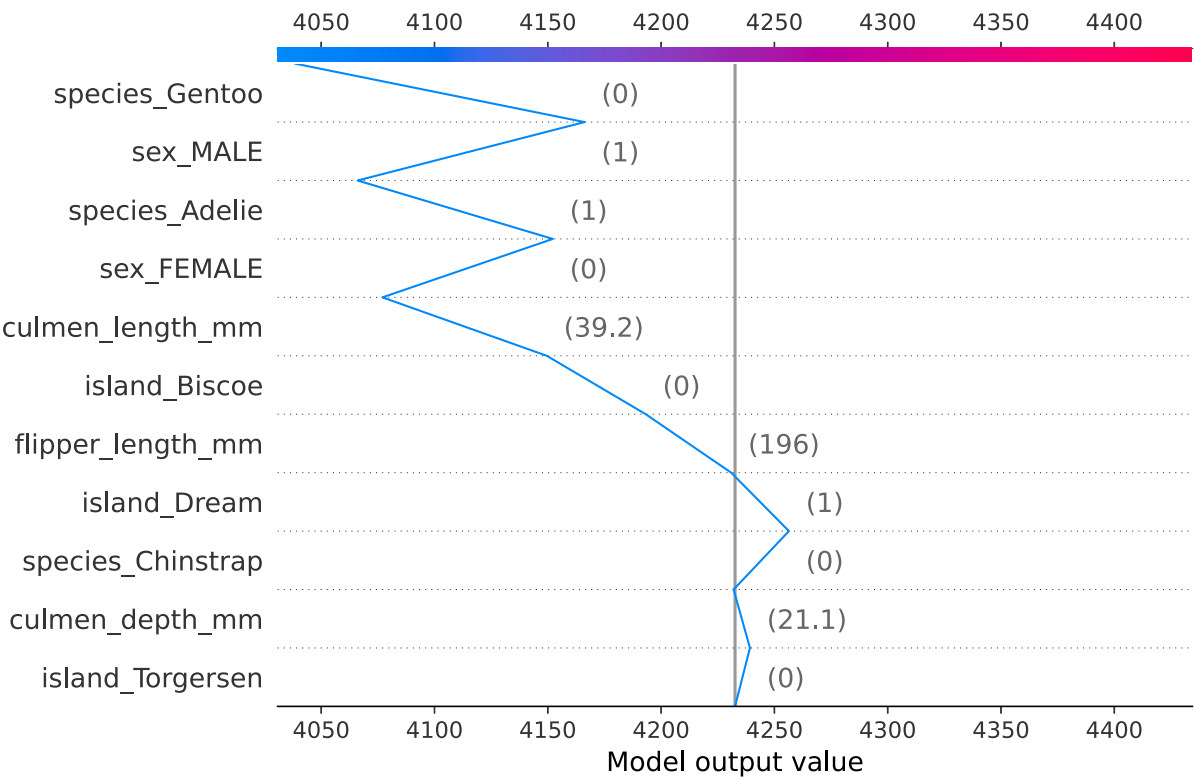
In [ ]: shap.decision\_plot(expected\_value, shap\_values2[33, :], X\_test.iloc[33, :])



In [ ]: p(7)

Out[ ]:


```
In [ ]: shap.decision_plot(expected_value, shap_values2[7, :], X_test.iloc[7, :])
```




```
In [ ]: p(27)
```

Out[ ]:

```
In [ ]: shap.decision_plot(expected_value, shap_values2[27, :], X_test.iloc[27, :])
```





In [ ]:

Loading [MathJax]/extensions/Safe.js